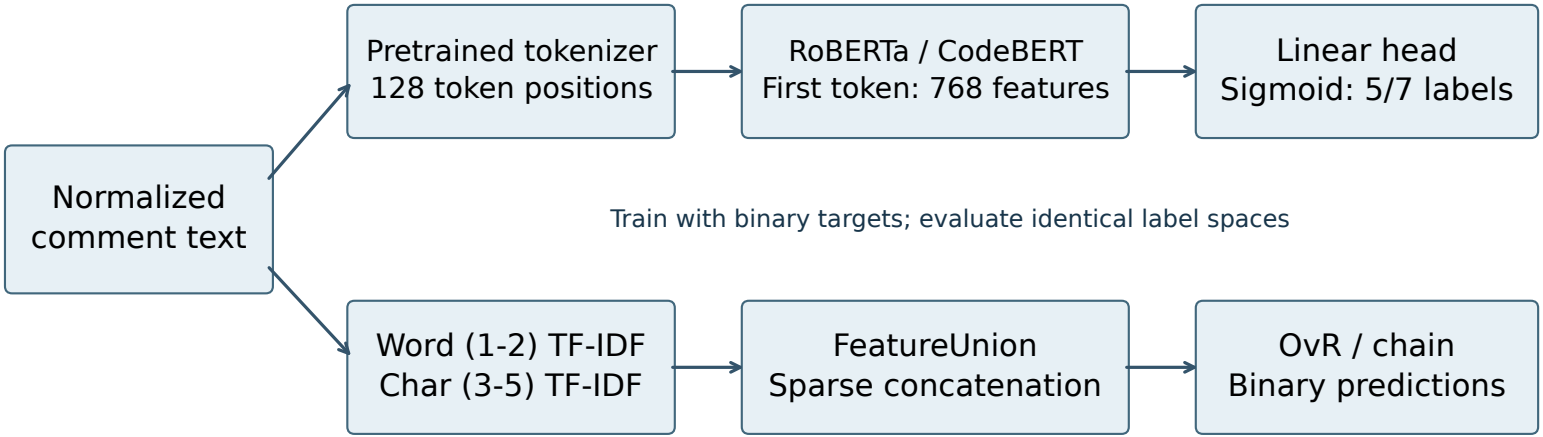
Normalized comment text
Pretrained tokenizer 128 token positions
RoBERTa / CodeBERT First token: 768 features
Linear head Sigmoid: 5/7 labels
Train with binary targets; evaluate identical label spaces
Word (1-2) TF-IDF Char (3-5) TF-IDF
FeatureUnion Sparse concatenation
OvR / chain Binary predictions